

It has been demonstrated that Multi-Agent (MA) problems have a wide variety of applications including surveying and exploration[1], [2], aiding construction management[3], water quality monitoring[4] and assisting in disaster evacuation[5]. Many techniques have been employed to these ends with some more common approaches including expert designed model based controllers[6]–[9] and Multi-Agent Reinforcement Learning (MARL)[10]–[14]. MARL has been the subject of much attention recently, with groups demonstrating its applicability in simulation and lab environments[10], [15], [16], as well as in controlled real-world settings[4], [12], [17]. However, the most commonly identified shortcoming of MARL is a lack of human interpretable explanations for why an agent is taking a certain action, although developing metrics for this is an active research area[18]–[20]. This work is crucially important for systems to operate in unstructured and uncontrolled real-world human environments where the ability to both prevent and diagnose errors are incredibly important.

Both traditionally designed model-based controllers and MARL controllers have attractive qualities and drawbacks. For example, model-based controllers are highly interpretable; they depend on a designer specifying a fixed set of rules and constraints who then derives controls laws which satisfy them. However, they suffer in terms of generality as a change in system dynamics often means rederiving the control law, and implementation for complicated tasks often require devising complicated decentralized logic to drive cooperation. MARL has almost the opposite set of issues where policies often lack interpretability but are highly general given proper training. Developing better tools to interpret the behavior of MARL algorithms is one of the major roadblocks to wide-scale real-world implementation. There are many ways to improve interpretability of MARL techniques, with various authors doing this by carefully designing network structures or employing post-hoc reasoning to rationalize decisions, among other methods[19], [20]. One promising approach from the past few years has been the use of Graph Neural Networks (GNNs) to improve model interpretability [12], [17], [21]. These models achieve better interpretability by representing MA systems as graphs where agents are nodes with information stored as relations between nodes, which proves to be a very interpretable perspective in MARL. To that end, the proposed research is the development and testing of interpretable GNNs for use in MA Systems acting in unstructured real-world environments.

The goal of the study is to produce a GNN based MARL framework which is capable of coordinating communication to allow agents to pathfind in partially observable real-world environments. The work extends current implementations of MARL to allow for better integration into everyday unstructured environments containing people. Phase 1 (6 months) will focus on development and in-lab verification of a set of interpretability metrics for MA systems. Phase 2 (6 months) focuses on deployment in unstructured but access-controlled real-world environments to ensure the algorithm behaves safely while giving interpretable rationale. Finally, phase 3 (1 year) will focus on scaling the study to real-world, human environments where uncontrolled interaction between agents and humans may occur. The required resources would be a collection of unmanned vehicles with hardware to enable decentralized communication and coordination and controlled space for experiments.

- [1] M. Kegeleirs, G. Grisetti, and M. Birattari, "Swarm SLAM: Challenges and Perspectives," *Front. Robot. AI*, vol. 8, p. 618268, Mar. 2021, doi: 10.3389/frobt.2021.618268.
- [2] D. He, D. Feng, H. Jia, and H. Liu, "Decentralized Exploration of a Structured Environment Based on Multi-agent Deep Reinforcement Learning," in *2020 IEEE 26th International Conference on Parallel and Distributed Systems (ICPADS)*, Hong Kong, Dec. 2020, pp. 172–179. doi: 10.1109/ICPADS51040.2020.00032.
- [3] L. Xiang, Y. Tan, G. Shen, and X. Jin, "Applications of multi-agent systems from the perspective of construction management: A literature review," *Eng. Constr. Archit. Manag.*, vol. 29, no. 9, pp. 3288–3310, Nov. 2022, doi: 10.1108/ECAM-01-2021-0038.
- [4] M. Rahmati, M. Nadeem, V. Sadhu, and D. Pompili, "UW-MARL: Multi-Agent Reinforcement Learning for Underwater Adaptive Sampling using Autonomous Vehicles," in *Proceedings of the International Conference on Underwater Networks & Systems*, Atlanta GA USA, Oct. 2019, pp. 1–5. doi: 10.1145/3366486.3366533.
- [5] A. R. Wagner, "Robot-Guided Evacuation as a Paradigm for Human-Robot Interaction Research," *Front. Robot. AI*, vol. 8, p. 701938, Jul. 2021, doi: 10.3389/frobt.2021.701938.
- [6] L. E. Beaver and A. A. Malikopoulos, "Constraint-Driven Optimal Control of Multi-Agent Systems: A Highway Platooning Case Study," *IEEE Control Syst. Lett.*, vol. 6, pp. 1754–1759, 2022, doi: 10.1109/LCSYS.2021.3133801.
- [7] L. E. Beaver and A. A. Malikopoulos, "Constraint-Driven Optimal Control for Emergent Swarming and Predator Avoidance." arXiv, Mar. 21, 2022. Accessed: Mar. 02, 2023. [Online]. Available: <http://arxiv.org/abs/2203.11057>
- [8] L. Wang, A. D. Ames, and M. Egerstedt, "Safety Barrier Certificates for Collisions-Free Multirobot Systems," *IEEE Trans. Robot.*, vol. 33, no. 3, pp. 661–674, Jun. 2017, doi: 10.1109/TRO.2017.2659727.
- [9] P. Varnai and D. V. Dimarogonas, "Multi-Agent Stochastic Control using Path Integral Policy Improvement," in *2022 American Control Conference (ACC)*, Atlanta, GA, USA, Jun. 2022, pp. 3406–3411. doi: 10.23919/ACC53348.2022.9867386.
- [10] P. D. Siedler, "Collaborative Auto-Curricula Multi-Agent Reinforcement Learning with Graph Neural Network Communication Layer for Open-ended Wildfire-Management Resource Distribution." arXiv, Apr. 24, 2022. Accessed: Mar. 02, 2023. [Online]. Available: <http://arxiv.org/abs/2204.11350>
- [11] E. S. Lee, L. Zhou, A. Ribeiro, and V. Kumar, "Graph Neural Networks for Decentralized Multi-Agent Perimeter Defense," *Front. Control Eng.*, vol. 4, p. 1104745, Jan. 2023, doi: 10.3389/fcteg.2023.1104745.
- [12] Z. Ye, X. Jiang, G. Song, and B. Yang, "Soft Hierarchical Graph Recurrent Networks for Many-Agent Partially Observable Environments." arXiv, Sep. 05, 2021. Accessed: Mar. 03, 2023. [Online]. Available: <http://arxiv.org/abs/2109.02032>
- [13] M. Dorigo, G. Theraulaz, and V. Trianni, "Reflections on the future of swarm robotics," *Sci. Robot.*, vol. 5, no. 49, p. eabe4385, Dec. 2020, doi: 10.1126/scirobotics.abe4385.
- [14] K. Zhang, Z. Yang, and T. Başar, "Decentralized multi-agent reinforcement learning with networked agents: recent advances," *Front. Inf. Technol. Electron. Eng.*, vol. 22, no. 6, pp. 802–814, Jun. 2021, doi: 10.1631/FITEE.1900661.

- [15] J. Blumenkamp, S. Morad, J. Gielis, Q. Li, and A. Prorok, "A Framework for Real-World Multi-Robot Systems Running Decentralized GNN-Based Policies," in *2022 International Conference on Robotics and Automation (ICRA)*, Philadelphia, PA, USA, May 2022, pp. 8772–8778. doi: 10.1109/ICRA46639.2022.9811744.
- [16] E. Seraj, "Learning Efficient Diverse Communication for Cooperative Heterogeneous Teaming," 2022.
- [17] G. Sartoretti, W. Paivine, Y. Shi, Y. Wu, and H. Choset, "Distributed Learning of Decentralized Control Policies for Articulated Mobile Robots," *IEEE Trans. Robot.*, vol. 35, no. 5, pp. 1109–1122, Oct. 2019, doi: 10.1109/TRO.2019.2922493.
- [18] F.-L. Fan, J. Xiong, M. Li, and G. Wang, "On Interpretability of Artificial Neural Networks: A Survey," *IEEE Trans. Radiat. Plasma Med. Sci.*, vol. 5, no. 6, pp. 741–760, Nov. 2021, doi: 10.1109/TRPMS.2021.3066428.
- [19] A. Alharin, T.-N. Doan, and M. Sartipi, "Reinforcement Learning Interpretation Methods: A Survey," *IEEE Access*, vol. 8, pp. 171058–171077, 2020, doi: 10.1109/ACCESS.2020.3023394.
- [20] C. Glanois *et al.*, "A Survey on Interpretable Reinforcement Learning." arXiv, Feb. 24, 2022. Accessed: Mar. 02, 2023. [Online]. Available: <http://arxiv.org/abs/2112.13112>
- [21] J. Sheng *et al.*, "Learning structured communication for multi-agent reinforcement learning," *Auton. Agents Multi-Agent Syst.*, vol. 36, no. 2, p. 50, Oct. 2022, doi: 10.1007/s10458-022-09580-8.